

Language Models

CE-324: Modern Information Retrieval

Sharif University of Technology

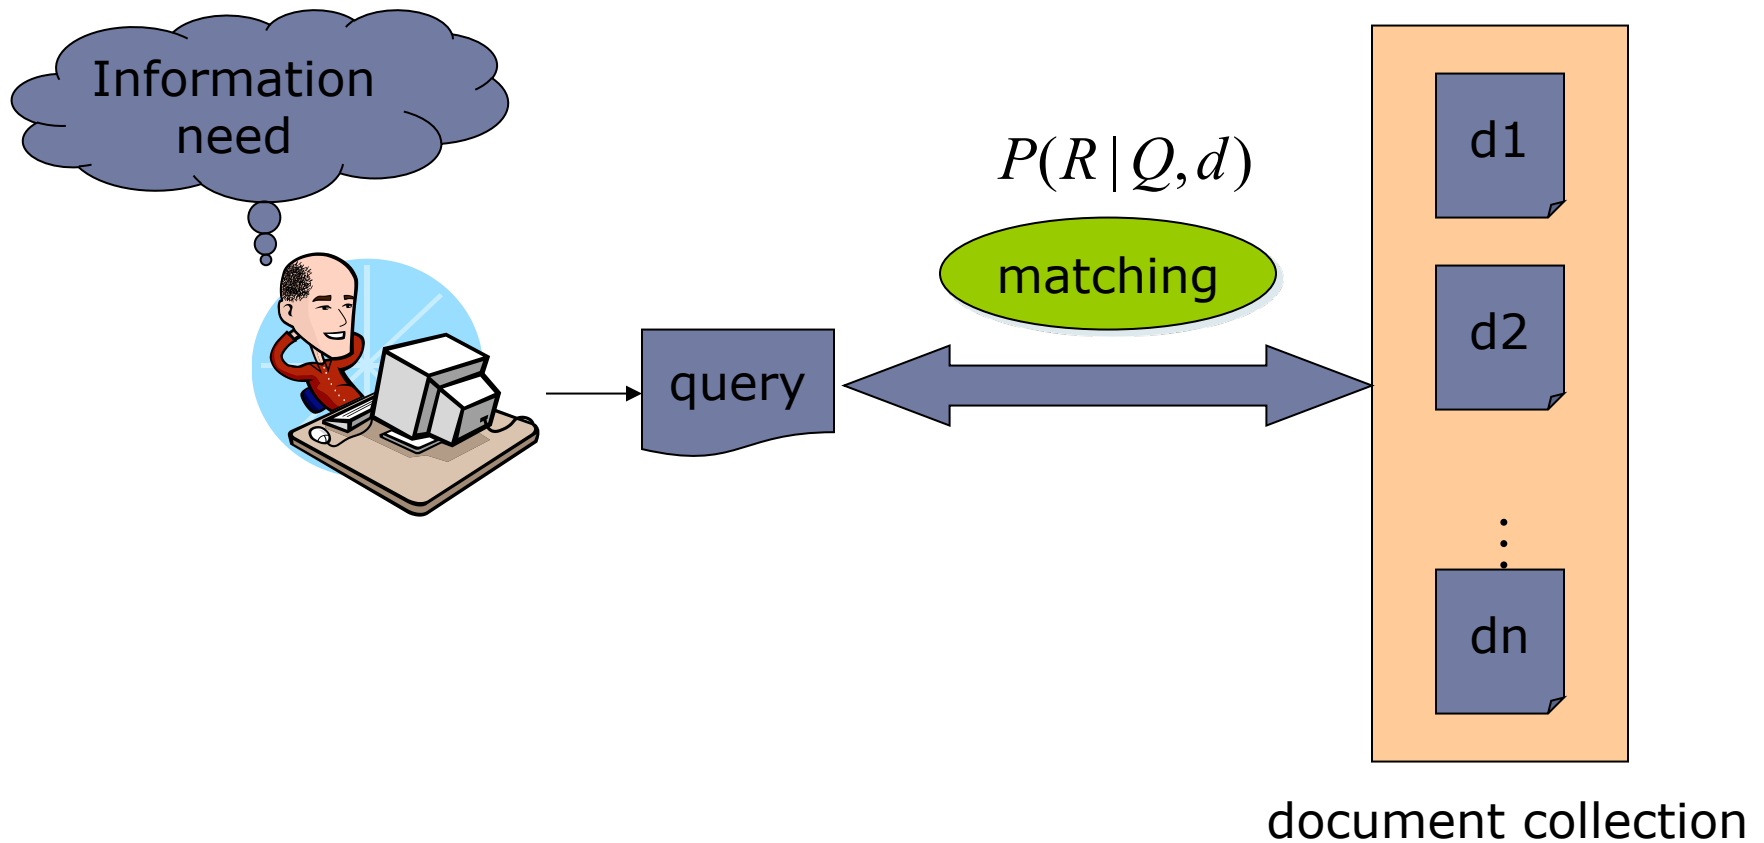
M. Soleymani

Spring 2020

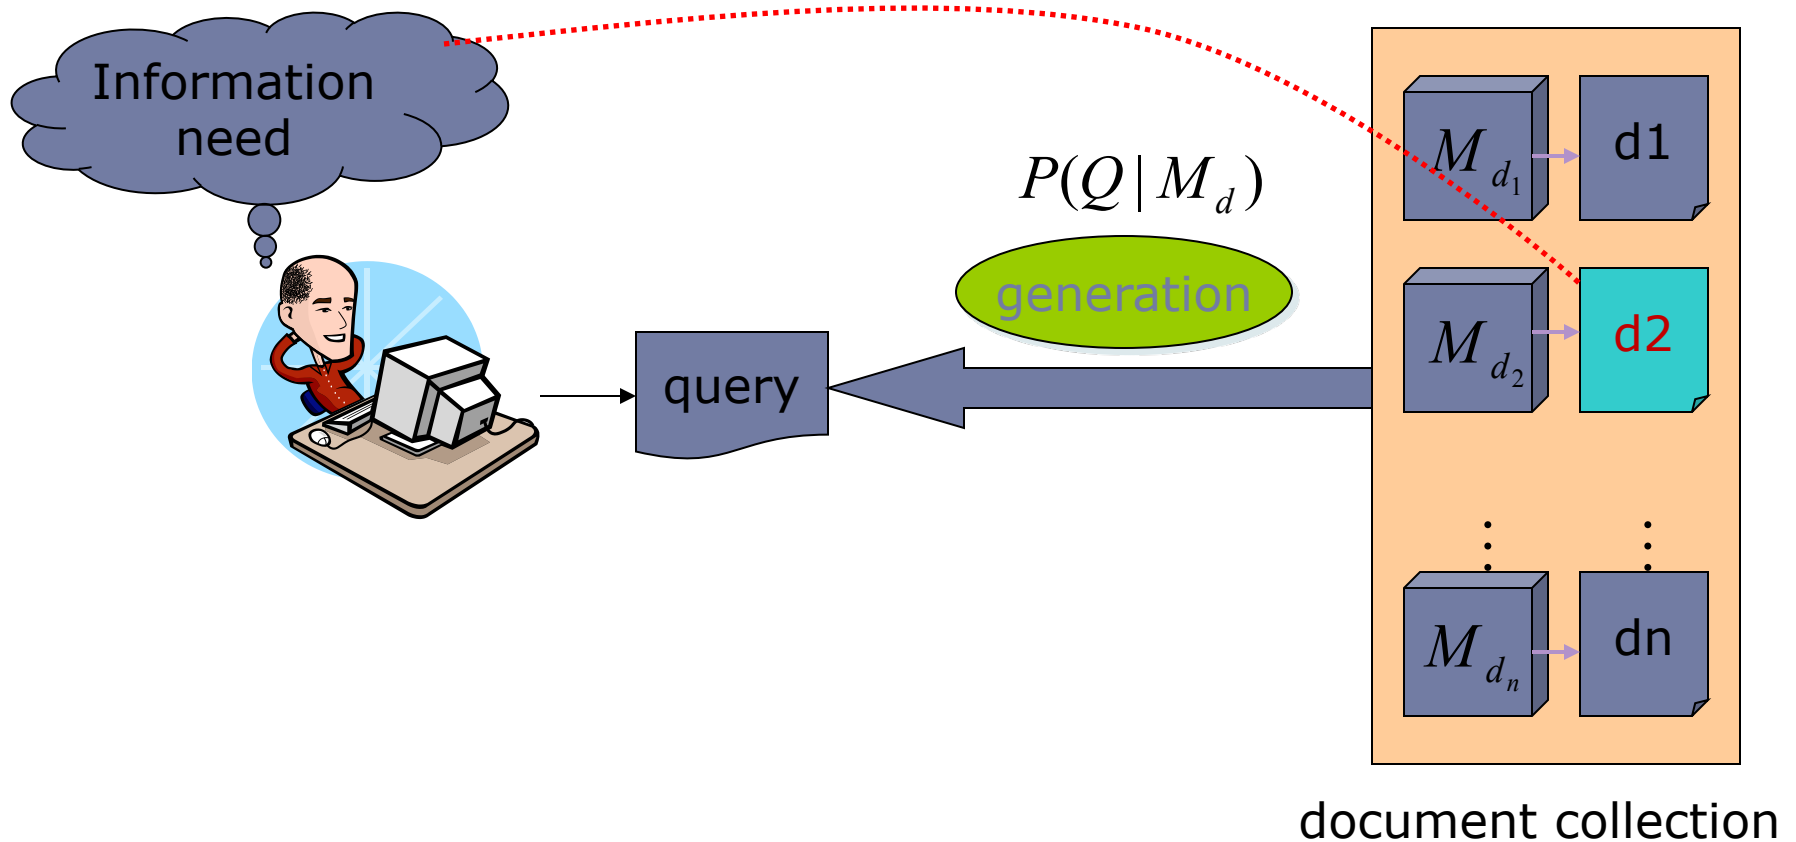
Most slides have been adapted from: Profs. Manning, Nayak & Raghavan (CS-276, Stanford)

Standard probabilistic IR: PRP

Ranking based on PRP



IR based on Language Model (LM)




Language models in IR

- ▶ Often, users have a reasonable idea of terms that are likely to occur in docs of interest
- ▶ They choose query terms that distinguish these docs from others in the collection
- ▶ LM approach assumes that docs and query are objects of the same type
 - ▶ Thus, assesses their match by importing the methods of language modeling

Formal language model

- ▶ Traditional generative model: generates strings
 - ▶ Finite state machines or regular grammars, etc.
- ▶ Example:



I wish

I wish I wish

I wish I wish I wish

I wish I wish I wish I wish

...

Stochastic language models

- ▶ Models *probability* of generating strings in the language (commonly all strings over alphabet Σ)

$$\sum_{s \in \Sigma^*} p(s) = 1$$

- ▶ Unigram model:
 - ▶ probabilistic finite automaton consisting of just a single node
 - ▶ with a single probability distribution over producing different terms $\sum_{t \in V} p(t) = 1$
 - ▶ also requires a probability of stopping in the finishing state

Example

Model M

the	0.2
a	0.1
information	0.01
retrieval	0.01
data	0.02
compute	0.03
...	

the information retrieval

0.2 0.01 0.01

multiply

$$P(s \mid M) \propto 0.00002$$

Stochastic language models

- Model *probability* of generating any string

Model M1	the	0.2
	a	0.1
	data	0.02
	information	0.01
	retrieval	0.01
	computing	0.005
	system	0.004

Model M2	the	0.15
	a	0.08
	management	0.05
	information	0.02
	database	0.02
	system	0.015
	mining	0.002

information	system
<hr/>	<hr/>
0.01	0.004
0.02	0.015

$$P(s|M_2) > P(s|M_1)$$

The fundamental problem of LMs


- ▶ Usually we don't know the model M
 - ▶ But have a sample of text representative of that model



- ▶ Estimate a language model from a sample doc




- ▶ Then compute the observation probability



Stochastic language models

- ▶ A statistical model for generating text
 - ▶ Probability distribution over strings in a given language



$$\begin{aligned} P(\text{red yellow red blue} \mid M) &= P(\text{red} \mid M) \times \\ &\quad P(\text{yellow} \mid \text{red}, M) \times \\ &\quad P(\text{red} \mid \text{red yellow}, M) \times \\ &\quad P(\text{blue} \mid \text{red yellow red}, M) \end{aligned}$$

Unigram and higher-order models

$$P(\text{red } \text{yellow } \text{red } \text{blue})$$

$$= P(\text{red}) P(\text{yellow} | \text{red}) P(\text{red} | \text{red } \text{yellow}) P(\text{blue} | \text{red } \text{yellow } \text{red})$$

► **Unigram** Language Models

$$P(\text{red}) P(\text{yellow}) P(\text{red}) P(\text{blue})$$



Easy.
Effective!


► **Bigram** (generally, n -gram) Language Models

$$P(\text{red}) P(\text{yellow} | \text{red}) P(\text{red} | \text{yellow}) P(\text{blue} | \text{red})$$

► Other Language Models

- Grammar-based models (PCFGs)
 - Probably not the first thing to try in IR

Unigram model



$$P(\text{red, yellow, red, blue}) = P(\text{red}) P(\text{yellow}) P(\text{red}) P(\text{blue})$$

Probabilistic language models in IR

- ▶ Treat each doc as the basis for a model
 - ▶ e.g., unigram sufficient statistics
- ▶ Rank doc d based on $P(d|q)$
 - ▶ $P(d|q) = P(q|d) \times P(d) / P(q)$
 - ▶ $P(q)$ is the same for all docs, so ignore
 - ▶ $P(d)$ [the prior] is often treated as the same for all d
 - But we could use criteria like authority, length, genre
 - ▶ $P(q|d)$ is the probability of q given d 's model
- ▶ Very general formal approach

Query likelihood language model

$$p(d|q) = \frac{p(q|d) \times p(d)}{p(q)}$$

$$\approx \frac{p(q|M_d) \times p(d)}{p(q)}$$

- ▶ Ranking formula

$$p(d)p(q | M_d)$$

Language models for IR

- ▶ Language Modeling Approaches

- ▶ Attempt to **model query generation process**

- ▶ Docs are ranked by **the probability that a query would be observed as a random sample from the doc model**

- ▶ Multinomial approach

$$P(q|M_d) = K_q \prod_{t \in V} P(t|M_d)^{\text{tf}_{t,q}}$$

$$K_q = \frac{L_q!}{\text{tf}_{1,q}! \times \cdots \times \text{tf}_{M,q}!}$$

Retrieval based on probabilistic LM

- ▶ Generation of queries as a random process
- ▶ Approach
 - ▶ Infer a language model for each doc.
 - ▶ Usually a unigram estimate of words is used
 - Some work on bigrams
 - ▶ Estimate the probability of generating the query according to each of these models.
 - ▶ Rank the docs according to these probabilities.

Query generation probability

- ▶ The probability of producing the query given the language model of doc d using MLE is:

$$\hat{p}(t|M_d) = \frac{tf_{t,d}}{L_d}$$

$$\hat{p}(q|M_d) \propto \prod_{t \in q} \hat{p}(t|M_d)^{tf_{t,q}}$$

Unigram assumption:
Given a particular language model, the query terms occur independently.

M_d : language model of document d

$tf_{t,d}$: raw tf of term t in document d

L_d : total number of tokens in document d

$tf_{t,q}$: raw tf of term t in query q

Insufficient data

- ▶ Zero probability

- ▶ May not wish to assign a probability of zero to a doc missing one or more of the query terms [gives conjunction semantics]

$$\hat{p}(t|M_d) = 0$$

- ▶ Poor estimation: occurring words may also be badly estimated
 - ▶ in particular, the probability of words occurring for example once in the doc is normally overestimated

Insufficient data: solution

- ▶ Zero probabilities spell disaster
 - ▶ We need to smooth probabilities
 - ▶ Discount nonzero probabilities
 - ▶ Give some probability mass to unseen things
- ▶ Smoothing: discounts non-zero probabilities and gives some probability mass to unseen words
- ▶ Many approaches to smoothing probability distributions to deal with this problem
 - ▶ i.e., adding 1, $1/2$ or α to counts, interpolation, and etc.

Collection statistics

- ▶ A non-occurring term is possible, but no more likely than would be expected by chance in the collection.

$$\text{If } t f_{t,d} = 0 \text{ then } \hat{p}(t|M_d) < \frac{c f_t}{T}$$

$c f_t$: raw count of term t in the collection

$c s = T$: raw collection size (total number of tokens in the collection)

$$\hat{p}(t|M_c) = \frac{c f_t}{T}$$

- ▶ Collection statistics ...
 - ▶ Are integral parts of the language model (as we will see).
 - ▶ Are not used heuristically as in many other approaches.
 - ▶ However there's some wiggle room for empirically set parameters

Bayesian smoothing

$$\hat{p}(t|d) = \frac{tf_{t,d} + \alpha \hat{p}(t|Mc)}{L_d + \alpha}$$

- ▶ For a word present in the doc:
 - ▶ combines a discounted MLE and a fraction of the estimate of its prevalence in the whole collection
- ▶ For words not present in a doc:
 - ▶ is just a fraction of the estimate of the prevalence of the word in the whole collection.

Linear interpolation: Mixture model

- ▶ **Linear interpolation:** Mixes the probability from the doc with the general collection frequency of the word. $0 \leq \lambda \leq 1$
- ▶ using a mixture between the doc multinomial and the collection multinomial distribution

$$\hat{p}(t|d) = \lambda \hat{p}(t|M_d) + (1 - \lambda) \hat{p}(t|Mc)$$

$$\hat{p}(t|d) = \lambda \frac{tf_{t,d}}{L_d} + (1 - \lambda) \frac{cf_t}{T}$$

- ▶ It works well in practice

Linear interpolation: Mixture model

- ▶ **Correctly setting λ is very important**
 - ▶ high value: “conjunctive-like” search— suitable for short queries
 - ▶ low value for long queries
 - ▶ Can tune λ to optimize performance
 - ▶ Perhaps make it dependent on doc size (cf. Dirichlet prior or Witten-Bell smoothing)

Basic mixture model: summary

► General formulation of the LM for IR

$$\hat{p}(q|d) = \prod_{t \in q} \lambda \hat{p}(t|M_d) + (1 - \lambda) \hat{p}(t|Mc)$$

general language model

individual-document model

- The user has a doc in mind, and generates the query from this doc.
- The equation represents the probability that the doc that the user had in mind was in fact this one.

Example

- ▶ Doc collection (2 docs)
 - ▶ d_1 : “Xerox reports a profit but **revenue** is **down**”
 - ▶ d_2 : “Lucent narrows quarter loss but **revenue** decreases further”
- ▶ Model: MLE unigram from docs; $\lambda = 1/2$
- ▶ Query: **revenue down**
 - ▶ $P(q|d_1) = [(1/8 + 2/16) / 2] \times [(1/8 + 1/16) / 2]$
 $= 1/8 \times 3/32 = 3/256$
 - ▶ $P(q|d_2) = [(1/8 + 2/16) / 2] \times [(0 + 1/16) / 2]$
 $= 1/8 \times 1/32 = 1/256$
- ▶ Ranking: $d_1 > d_2$

Ponte and croft experiments

▶ Data

- ▶ TREC topics 202-250 on TREC disks 2 and 3
 - ▶ Natural language queries consisting of one sentence each
- ▶ TREC topics 51-100 on TREC disk 3 using the concept fields
 - ▶ Lists of good terms

<num>Number: 054

<dom>Domain: International Economics

<title>Topic: Satellite Launch Contracts

<desc>Description:

... </desc>

<con>Concept(s):

1. Contract, agreement
2. Launch vehicle, rocket, payload, satellite
3. Launch services, ... </con>

Precision/recall results 202-250

Rec.	Precision			
	tf-idf	LM	%chg	
0.0	0.7439	0.7590	+2.0	
0.1	0.4521	0.4910	+8.6	
0.2	0.3514	0.4045	+15.1	*
0.3	0.2761	0.3342	+21.0	*
0.4	0.2093	0.2572	+22.9	*
0.5	0.1558	0.2061	+32.3	*
0.6	0.1024	0.1405	+37.1	*
0.7	0.0451	0.0760	+68.7	*
0.8	0.0160	0.0432	+169.6	*
0.9	0.0033	0.0063	+89.3	
1.0	0.0028	0.0050	+76.9	
Ave	0.1868	0.2233	+19.55	*

LM vs. probabilistic model for IR (PRP)

- ▶ Main difference: whether “Relevance” figures explicitly in the model or not
 - ▶ LM approach attempts to do away with modeling relevance
- ▶ LM approach assumes that docs and queries are of the same type
- ▶ Computationally tractable, intuitively appealing

LM vs. probabilistic model for IR

- ▶ Problems of basic LM approach
 - ▶ Assumption of equivalence between doc and information problem representation is unrealistic
 - ▶ Very simple models of language
 - ▶ Relevance feedback is difficult to integrate
 - ▶ user preferences, and other general issues of relevance
 - ▶ Can't easily accommodate phrases, passages, Boolean operators
- ▶ Recent work has shown the LM approach to be very effective in retrieval experiments, beating tf-idf and BM25 weights

Translation model (Berger and Lafferty)

- ▶ Basic LMs do not address issues of synonymy.
 - ▶ Or any deviation in expression of information need from language of docs
- ▶ A translation model: generate query words not in doc via “translation” to synonyms etc.
 - ▶ Or to do cross-language IR, or multimedia IR

$$P(q | M_d) = \prod_{t \in q} \sum_{v \in V} P(v | M_d) \times T(t | v)$$

Basic LM Translation

- ▶ Need to learn a translation model (using a dictionary or via statistical machine translation)

Language models: summary

- ▶ Novel way of looking at IR problem based on probabilistic language modeling
 - ▶ Conceptually simple and explanatory
 - ▶ Formal mathematical model
 - ▶ Natural use of collection statistics, not heuristics (almost...)
- ▶ Effective retrieval and can be improved to the extent that the following conditions can be met
 - ▶ accurate representations of the data
 - ▶ users have some sense of term distribution
 - ▶ we get more sophisticated with translation model

Comparison with vector space

- ▶ There's some relation to traditional tf.idf models:
 - ▶ (unscaled) term frequency is directly in model
 - ▶ probabilities do length normalization of term frequencies
 - ▶ effect of doing a mixture with overall collection frequencies is a little like idf:
 - ▶ terms rare in the general collection but common in some documents will have a greater influence on the ranking

Comparison with vector space

- ▶ Similar in some ways
 - ▶ Term weights based on their frequency
 - ▶ Terms often used as if they were independent
 - ▶ Inverse document/collection frequency used
 - ▶ Some form of length normalization useful
- ▶ Different in others
 - ▶ Based on probability rather than similarity
 - ▶ Intuitions are probabilistic rather than geometric
 - ▶ Details of use of document length and term, document, and collection frequency differ

Resources

IIR, Chapter 12.

The Lemur Toolkit for Language Modeling and Information Retrieval. [CMU/Umass LM and IR system in C(++)] <http://www-2.cs.cmu.edu/~lemur/>