

# Text classification I (Naïve Bayes)

CE-324: Modern Information Retrieval

Sharif University of Technology

M. Soleymani

Spring 2020

Most slides have been adapted from: Profs. Manning, Nayak & Raghavan (CS-276, Stanford)

# Outline

---

- ▶ Text classification
  - ▶ definition
  - ▶ relevance to information retrieval
- ▶ Naïve Bayes classifier

# Formal definition of text classification

---

- ▶ Document space  $X$ 
  - ▶ Docs are represented in this (typically high-dimensional) space
- ▶ Set of classes  $\mathcal{C} = \{c_1, \dots, c_K\}$ 
  - ▶ Example:  $\mathcal{C} = \{\text{spam}, \text{non-spam}\}$
- ▶ Training set: a set of labeled docs. Each labeled doc  $\langle d, c \rangle \in X \times \mathcal{C}$
  
- ▶ Using a learning method, we find a classifier  $\gamma(\cdot)$  that maps docs to classes:  $\gamma: X \rightarrow \mathcal{C}$

# Examples of using classification in IR systems

---

- ▶ Language identification (classes: English vs. French etc.)
- ▶ Automatic detection of spam pages (spam vs. non-spam)
- ▶ Automatic detection of secure pages for safe search
- ▶ Topic-specific or vertical search – restrict search to a “vertical” like “related to health” (relevant to vertical vs. not)
- ▶ Sentiment detection: is a movie or product review positive or negative (positive vs. negative)
- ▶ Exercise: Find examples of uses of text classification in IR

# Standing queries

---

- ▶ The path from IR to text classification:
  - ▶ You have an information need to monitor, say:
    - ▶ Unrest in the Niger delta region
  - ▶ You want to rerun an appropriate query periodically to find new news items on this topic
  - ▶ You will be sent new documents that are found
    - ▶ I.e., it's not ranking but classification (relevant vs. not relevant)
- ▶ Such queries are called **standing queries**
  - ▶ Long used by “information professionals”
  - ▶ A modern mass instantiation is **Google Alerts**

**From:** Google Alerts  
**Subject:** Google Alert - stanford -neuro-linguistic nlp OR "Natural Language Processing" OR parser OR tagger OR ner OR "named entity" OR segmenter OR classifier OR dependencies OR "core nlp" OR corenlp OR phrasal  
**Date:** May 7, 2012 8:54:53 PM PDT  
**To:** Christopher Manning

---

<b>Web</b>	3 new results for stanford -neuro-linguistic nlp OR "Natural Language Processing" OR parser OR tagger OR ner OR "named entity" OR segmenter OR classifier OR dependencies OR "core nlp" OR corenlp OR phrasal
------------	---

[Twitter / Stanford NLP Group: @Robertoross If you only n ...](#)

@Robertoross If you only need tokenization, java -mx2m edu.stanford.nlp.process.PTBTOKENIZER file.txt runs in 2MB on a whole file for me.... 9:41 PM Apr 28th ...  
[twitter.com/stanfordnlp/status/196459102770171905](https://twitter.com/stanfordnlp/status/196459102770171905)

[\[Java\] LexicalizedParser lp = LexicalizedParser.loadModel\("edu ...](#)

loadModel("edu/stanford/nlp/models/lexparser/englishPCFG.ser.gz"); String[] sent = { "This", "is", "an", "easy", "sentence", "." }; Tree parse = lp.apply(Arrays.  
[pastebin.com/az14R9nd](http://pastebin.com/az14R9nd)

[More Problems with Statistical NLP || kuro5hin.org](#)

Tags: nlp, ai, coursera, stanford, nlp-class, cky, nltk, reinventing the wheel, ... Programming Assignment 6 for Stanford's nlp-class is to implement a CKY parser .  
[www.kuro5hin.org/story/2012/5/5/11011/68221](http://www.kuro5hin.org/story/2012/5/5/11011/68221)

---

Tip: Use quotes ("like this") around a set of words in your query to match them exactly. [Learn more.](#)

[Delete](#) this alert.  
[Create](#) another alert.  
[Manage](#) your alerts.

# Spam filtering

## Another text classification task

---

From: "" <takworld@hotmail.com>

Subject: real estate is the only way... gem oalvgkay

Anyone can buy real estate with no money down

Stop paying rent TODAY !

There is no need to spend hundreds or even thousands for similar courses

I am 22 years old and I have already purchased 6 properties using the methods outlined in this truly INCREDIBLE ebook.

Change your life NOW !

=====

Click Below to order:

<http://www.wholesaledaily.com/sales/nmd.htm>

=====

# Categorization/Classification

---

## ▶ Given:

- ▶ A representation of a document  $d$ 
  - ▶ Issue: how to represent text documents.
  - ▶ Usually some type of high-dimensional space – bag of words

- ▶ A fixed set of classes:

$$\mathcal{C} = \{c_1, c_2, \dots, c_J\}$$

## ▶ Determine:

- ▶ The category of  $d$ :  $\gamma(d) \in \mathcal{C}$ 
  - ▶  $\gamma(d)$  is a classification function
- ▶ We want to build classification functions (“classifiers”).



# Classification Methods (1)

---

- Manual classification
  - Used by the original Yahoo! Directory
  - Looksmart, about.com, ODP, PubMed
- Accurate when job is done by experts
- Consistent when the problem size and team is small
- Difficult and expensive to scale
  - Means we need automatic classification methods for big problems

# Classification Methods (2)

---

## ▶ Hand-coded rule-based classifiers

- ▶ One technique used by news agencies, intelligence agencies, etc.
- ▶ Widely deployed in government and enterprise
- ▶ Vendors provide “IDE” for writing such rules

## ▶ Issues:

- ▶ Commercial systems have complex query languages
- ▶ Accuracy can be high if a rule has been carefully refined over time by a subject expert
- ▶ Building and maintaining these rules is expensive

# Classification Methods (3):

## Supervised learning

---

### ▶ Given:

- ▶ A document  $d$
- ▶ A fixed set of classes:  
 $C = \{c_1, c_2, \dots, c_J\}$
- ▶ A training set  $D$  of documents each with a label in  $C$

### ▶ Determine:

- ▶ A learning method or algorithm which will enable us to learn a classifier  $\gamma$
- ▶ For a test document  $d$ , we assign it the class  
 $\gamma(d) \in C$

# Bayes classifier

---

- ▶ Bayesian classifier is a probabilistic classifier:

$$c = \operatorname{argmax}_k P(C_k | d)$$
$$c = \operatorname{argmax}_k P(d | C_k) P(C_k)$$

- ▶  $d = \langle t_1, \dots, t_{L_d} \rangle$
- ▶ There are too many parameters  $P(\langle t_1, \dots, t_{L_d} \rangle | C_k)$ 
  - ▶ One for each unique combination of a class and a sequence of words.
  - ▶ We would need a very, very large number of training examples to estimate that many parameters.

# Naïve bayes assumption

---

- ▶ Naïve bayes assumption:

$$P(d|C_k) = P(\langle t_1, \dots, t_{L_d} \rangle | C_k) = \prod_{i=1}^{L_d} P(t_i | C_k)$$

- ▶  $L_d$ : length of doc  $d$  (number of tokens)
- ▶  $P(t_i | C_k)$ : probability of term  $t_i$  occurring in a doc of class  $C_k$
- ▶  $P(C_k)$ : prior probability of class  $C_k$ .

# Naïve bayes assumption

---

- ▶ Naïve bayes assumption:

$$P(d|C_k) = P(\langle t_1, \dots, t_{L_d} \rangle | C_k) = \prod_{i=1}^{L_d} P(t_i | C_k)$$

- ▶  $L_d$ : length of doc  $d$  (number of tokens)
  - ▶  $P(t_i | C_k)$ : probability of term  $t_i$  occurring in a doc of class  $C_k$
  - ▶  $P(C_k)$ : prior probability of class  $C_k$ .
- 
- ▶ Equivalent to (language model view):

$$P(d|C_k) = \prod_{i=1}^{|V|} P(t_i | C_k)^{tf_{t_i,d}}$$

# Naive Bayes classifier

---

- ▶ Since  $\log$  is a monotonic function, the class with the highest score does not change.

$$c = \operatorname{argmax}_k P(d|C_k)P(C_k) = \operatorname{argmax}_k P(C_k) \prod_{i=1}^{L_d} P(t_i|C_k)$$

$$c = \operatorname{argmax}_k \log P(C_k) + \sum_{i=1}^{L_d} \log P(t_i|C_k)$$

$\log P(t_i|C_k)$ : a weight that indicates how good an indicator  $t_i$  is for  $C_k$

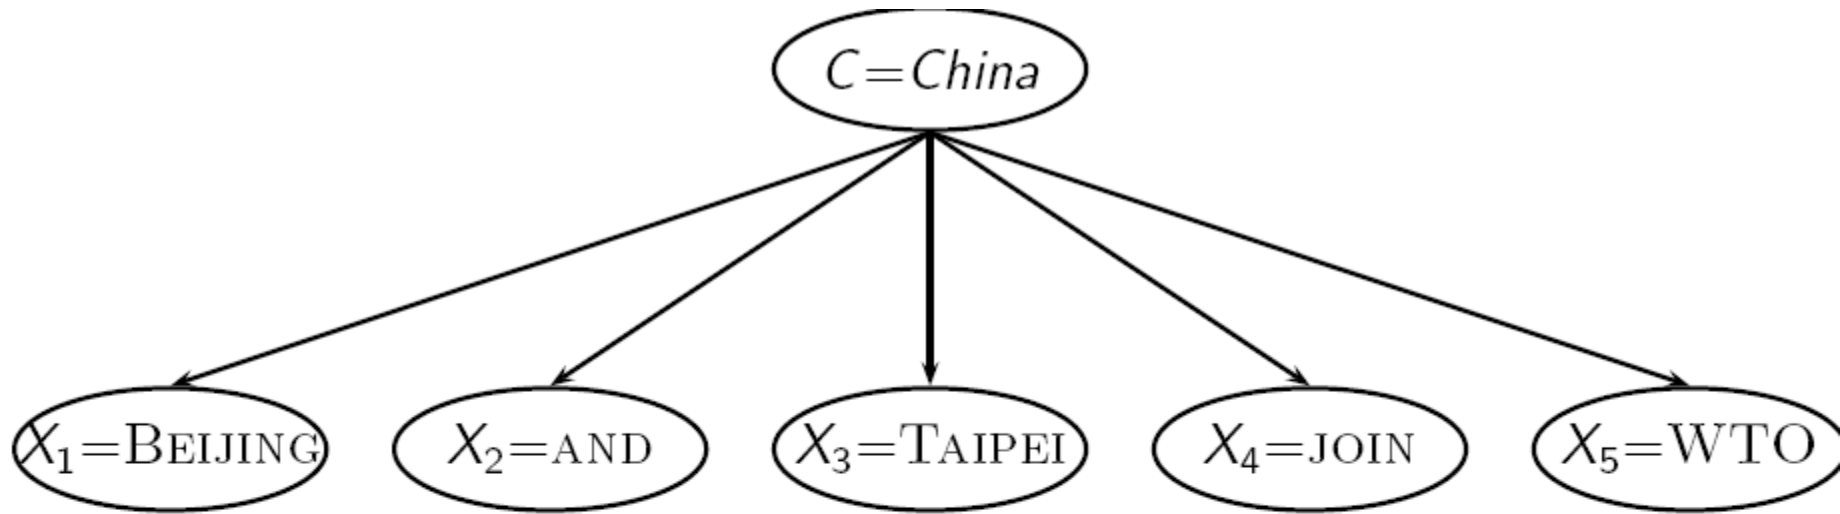
# Estimating parameters

---

- ▶ Estimate  $\hat{P}(C_k)$  and  $\hat{P}(t_i|C_k)$  from training data
  - ▶  $N_k$ : number of docs in class  $C_k$
  - ▶  $T_{i,k}$ : number of occurrence of  $t_i$  in training docs from class  $C_k$  (includes multiple occurrences)
- ▶  $\hat{P}(C_k) = \frac{N_k}{N}$
- ▶  $\hat{P}(t_i|C_k) = \frac{T_{i,k}}{\sum_{j=1}^M T_{j,k}}$



# Problem with estimates: Zeros



$$P(China|d) \propto P(China) \cdot P(\text{BEIJING}|China) \cdot P(\text{AND}|China) \\ \cdot P(\text{TAIPEI}|China) \cdot P(\text{JOIN}|China) \cdot P(\text{WTO}|China)$$

*d: BEIJING AND TAIPEI JOIN WTO*

$$P(WTO|China) = 0$$

# Problem with estimates: Zeros

---

- ▶ For doc  $d$  containing a term  $t$  that does not occur in any doc of a class  $c \Rightarrow \hat{P}(c|d) = 0$ 
  - ▶ Thus  $d$  cannot be assigned to class  $c$

- ▶ We use

$$\hat{P}(t|c) = \frac{T_{t,c} + 1}{(\sum_{t' \in V} T_{t',c}) + |V|}$$

Instead of

$$\hat{P}(t|c) = \frac{T_{t,c}}{\sum_{t' \in V} T_{t',c}}$$

# Naïve Bayes: summary

---

- ▶ Estimate parameters from the training corpus using add-one smoothing
- ▶ For a new doc  $d = t_1, \dots, t_{L_d}$ , for each class, compute  $\log P(C_k) + \sum_{i=1}^{L_d} \log P(t_i|C_k)$
- ▶ Assign doc  $d$  to the class with the largest score

# Naïve Bayes: example

---

	docID	words in document	in $c = \textit{China}$ ?
training set	1	Chinese Beijing Chinese	yes
	2	Chinese Chinese Shanghai	yes
	3	Chinese Macao	yes
	4	Tokyo Japan Chinese	no
test set	5	Chinese Chinese Chinese Tokyo Japan	?

- ▶ Training phase:
  - ▶ Estimate parameters of Naive Bayes classifier
- ▶ Test phase
  - ▶ Classifying the test doc

# Naïve Bayes: example

$C = \text{China}$

## ▶ Estimating parameters

$$\square \hat{P}(C) = \frac{3}{4}, \hat{P}(\bar{C}) = \frac{1}{4}$$

$$\square \hat{P}(\text{CHINESE}|C) = \frac{5+1}{8+6} = \frac{6}{14} \quad \hat{P}(\text{CHINESE}|\bar{C}) = \frac{1+1}{3+6} = \frac{2}{9}$$


$$\square \hat{P}(\text{TOKYO}|C) = \frac{0+1}{8+6} = \frac{1}{14} \quad \hat{P}(\text{TOKYO}|\bar{C}) = \frac{1+1}{3+6} = \frac{2}{9}$$

$$\square \hat{P}(\text{JAPAN}|C) = \frac{0+1}{8+6} = \frac{1}{14} \quad \hat{P}(\text{JAPAN}|\bar{C}) = \frac{1+1}{3+6} = \frac{2}{9}$$

## ▶ Classifying the test doc:

$$\triangleright \hat{P}(C|d) \propto \frac{3}{4} \times \left(\frac{6}{14}\right)^3 \times \frac{1}{14} \times \frac{1}{14} \approx 0.0003$$

$$\triangleright \hat{P}(\bar{C}|d) \propto \frac{1}{4} \times \left(\frac{2}{9}\right)^3 \times \frac{2}{9} \times \frac{2}{9} \approx 0.0001$$



$\hat{c} = C$

# Naïve Bayes: training

---

TRAINMULTINOMIALNB( $\mathbb{C}, \mathbb{D}$ )

1  $V \leftarrow \text{EXTRACTVOCABULARY}(\mathbb{D})$

2  $N \leftarrow \text{COUNTDOCS}(\mathbb{D})$

3 **for each**  $c \in \mathbb{C}$

4 **do**  $N_c \leftarrow \text{COUNTDOCSINCLASS}(\mathbb{D}, c)$

5  $\text{prior}[c] \leftarrow N_c / N$

6  $\text{text}_c \leftarrow \text{CONCATENATETEXTOFALLDOCSINCLASS}(\mathbb{D}, c)$

7 **for each**  $t \in V$

8 **do**  $T_{ct} \leftarrow \text{COUNTTOKENSOFTERM}(\text{text}_c, t)$

9 **for each**  $t \in V$

10 **do**  $\text{condprob}[t][c] \leftarrow \frac{T_{ct}+1}{\sum_{t'} (T_{ct'}+1)}$

11 **return**  $V, \text{prior}, \text{condprob}$

# Naïve Bayes: test

---

```
APPLYMULTINOMIALNB( $\mathbb{C}$ ,  $V$ ,  $prior$ ,  $condprob$ ,  $d$ )
1   $W \leftarrow \text{EXTRACTTOKENSFROMDOC}(V, d)$ 
2  for each  $c \in \mathbb{C}$ 
3  do  $score[c] \leftarrow \log prior[c]$ 
4      for each  $t \in W$ 
5      do  $score[c] + = \log condprob[t][c]$ 
6  return  $\arg \max_{c \in \mathbb{C}} score[c]$ 
```

# Time complexity of Naive Bayes

mode	time complexity	
training	$\Theta( \mathbb{D} L_{ave} +  \mathbb{C}  V )$	Generally: $ \mathbb{C}  V  <  D L_{ave}$
testing	$\Theta(L_a +  \mathbb{C} M_a) = \Theta( \mathbb{C} M_a)$	

- ▶  $D$ : training set,  $V$ : vocabulary,  $\mathbb{C}$ : set of classes
- ▶  $L_{ave}$ : average length of a training doc
- ▶  $L_a$ : length of the test doc
- ▶  $M_a$ : number of distinct terms in the test doc
- ▶ Thus: Naive Bayes is **linear** in the size of the training set (**training**) and the test doc (**testing**).
  - ▶ This is optimal time.



# Why does Naive Bayes work?

---

- ▶ The independence assumptions do not really hold of docs written in natural language.
- ▶ Naive Bayes can work well even though these assumptions are badly violated.
- ▶ Classification is about predicting the correct class and not about accurately estimating probabilities.
  - ▶ Naive Bayes is terrible for correct estimation ...
  - ▶ but it often performs well at choosing the correct class.

# Naive Bayes is not so naive

---

- ▶ Naive Bayes has won some bakeoffs (e.g., KDD-CUP 97)
- ▶ A good dependable baseline for text classification (but not the best)
  - ▶ Optimal if independence assumptions hold (never true for text, but true for some domains)
  - ▶ More robust to non-relevant features than some more complex learning methods
  - ▶ More robust to concept drift (changing of definition of class over time) than some more complex learning methods
- ▶ Very fast
- ▶ Low storage requirements

# Resources

---

- ▶ Chapter 13 of IIR